

CS 109B/STAT 121B/AC 209B/CSCI E-109B: Homework 6

Neural Networks - CNNs and RNNs

Harvard University

Spring 2018

Instructors: Pavlos Protopapas and Mark Glickman

INSTRUCTIONS

- To submit your assignment follow the instructions given in canvas.
- Restart the kernel and run the whole notebook again before you submit.
- Do not include your name(s) in the notebook if you are submitting as a group.
- If you submit individually and you have worked with someone, please include the name of your [one] partner below.

Your partner's name (if you submit separately):

Enrollment Status (109B, 121B, 209B, or E109B): 209B

Problem 1: Convolutional Neural Network Basics (10 pts)

In convolutional neural networks, a convolution is a multiplicative operation on a local region of values. Convolutional layers have been very useful in image classification, as it allows the network to retain local spatial information for feature extraction.

Part A: Understanding Convolutional Operations

For the following 2D matrix:

1	2	2
3	1	2
4	1	0

Use the following 2x2 kernel to perform a 2D convolution on the matrix:

2	1
1	2

1. Compute this operation by hand assuming a) valid, b) same, and c) full border modes. Please indicate what the resulting matrix shape is compared to the original shape.

Answers

a) Valid border - the shape of the resulting matrix shrinks by (kernel width - 1 = 2 - 1 = 1), compared to the original shape. Therefore, the original matrix shrinks from 3x3 to 2x2.

$$\text{output} = \begin{bmatrix} 9 & 11 \\ 13 & 5 \end{bmatrix} \text{ (shape: 2 by 2)}$$

b) Same border - Pad zeros to the input so that the resulting convolution feature map is of the same size as the original input (3x3).

$$\text{If zeros are padded on the right and bottom edges, input with padding} = \begin{bmatrix} 1 & 2 & 2 & 0 \\ 3 & 1 & 2 & 0 \\ 4 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} \text{ (shape: 4 by 4)}$$

$$\text{output} = \begin{bmatrix} 9 & 11 & 6 \\ 13 & 5 & 4 \\ 9 & 2 & 0 \end{bmatrix} \text{ (shape: 3 by 3)}$$

$$\text{If zeros are padded on the left and top edges, input with padding} = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 2 \\ 0 & 3 & 1 & 2 \\ 0 & 4 & 1 & 0 \end{bmatrix} \text{ (shape: 4 by 4)}$$

$$\text{output} = \begin{bmatrix} 2 & 5 & 6 \\ 7 & 9 & 11 \\ 11 & 13 & 5 \end{bmatrix} \text{ (shape: 3 by 3)}$$

c) Full border - Pad zeros so that each pixel in the original input is visited the same number of times (which is 4 in this case).

$$\text{input with padding} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 2 & 0 \\ 0 & 3 & 1 & 2 & 0 \\ 0 & 4 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \text{ (shape: 5 by 5)}$$
$$\text{output} = \begin{bmatrix} 2 & 5 & 6 & 2 \\ 7 & 9 & 11 & 6 \\ 11 & 13 & 5 & 4 \\ 4 & 9 & 2 & 0 \end{bmatrix} \text{ (shape: 4 by 4)}$$

Part B: Understanding Pooling Operations

Pooling operations are used in convolutional neural networks to reduce the dimensionality of the feature maps and overall network complexity. Two main types of pooling are used in CNNs: AveragePooling and MaxPooling.

1. Using the matrix below, write the output of the AveragePooling and MaxPooling operations with a pool size of 2x2 and stride 2x2. Then, write the outputs for the same operations, except with a stride size of 1.

1	2	2	4
3	1	2	1
4	1	0	2
5	2	2	1

Answers

a) AveragePooling - stride 2x2

$$\begin{bmatrix} 1.75 & 2.25 \\ 3.0 & 1.25 \end{bmatrix}$$

b) MaxPooling - stride 2x2

$$\begin{bmatrix} 3 & 4 \\ 5 & 2 \end{bmatrix}$$

c) AveragePooling - stride 1x1

$$\begin{bmatrix} 1.75 & 1.75 & 2.25 \\ 2.25 & 1.0 & 1.25 \\ 3.0 & 1.25 & 1.25 \end{bmatrix}$$

d) MaxPooling - stride 1x1

$$\begin{bmatrix} 3 & 2 & 4 \\ 4 & 2 & 2 \\ 5 & 2 & 2 \end{bmatrix}$$

Part C: Puppy Example

Consider the following image of a dog, which you will find in `dog.jpg` :



Load the image as a 2D Numpy array. Normalize the image by the following operation so that values fall within [-0.5, 0.5].

Perform the following steps for four images:

1. Randomly generate a 3x3 kernel.

2. Use this kernel and convolve over the image with same border mode (with `scipy.signal.convolve2d` (<https://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.convolve2d.html>)).

3. In the resulting image, set all pixel values less than zero to zero (using `np.clip()`). In other words:

```
if x < 0:
    x = 0
else:
    x = x
```

(This is the ReLU activation function.)

4. Plot the image.


Take a moment to examine the convolved images. You should see that certain features in the puppy are accentuated, while others are de-emphasized. Now consider the effect of performing additional convolution operations on these filtered images and how they relate to additional layers in a neural network.

```
In [1]: 1 import numpy as np
2 import scipy as sp
3 from scipy import signal
4 from imageio import imread
5 from copy import deepcopy
6 import matplotlib.pyplot as plt
7 import matplotlib.image as mpimg
8
9 import os
10 import keras
11 from keras.models import Sequential, load_model
12 from keras.layers import Dense, Dropout, Conv2D, Flatten, MaxPooling2D, Activation, Embedding, LSTM
13 from keras.optimizers import SGD
14 from keras import regularizers
15 from keras.datasets import cifar10
16
17 from keras.models import Model
18
19 %matplotlib inline
```

```
/usr/share/anaconda3/lib/python3.6/site-packages/h5py/_init_.py:36: FutureWarning: Conversion of the second argument of issubdtype
from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [2]: 1 def normalize(X, a, b):
2     X_norm = deepcopy(X)
3     if X.max() - X.min() != 0:
4         X_norm = (b-a) * (X - X.min()) / (X.max() - X.min()) + a
5     return X_norm
```


```
In [3]: 1 np.random.seed(900)
2 # read in dog image as np array
3 img = np.array(imread("dog.jpg"))
4
5 fig, ax = plt.subplots(1, 4, figsize=(25, 5))
6
7 # generate 4 different random kernels to see different affects
8 for i in range(4):
9
10     # randomly generate a 3x3 kernel
11     rand_kernel = np.random.randn(3,3)
12
13     # normalize the values to fall within [-0.5, 0.5]
14     img_norm = normalize(img, -0.5, 0.5)
15
16     # apply 2D convolution
17     img_conv = signal.convolve2d(img_norm, rand_kernel, mode='same')
18
19     # apply relu
20     img_relu = np.clip(img_conv, 0, None)
21     ax[i].imshow(img_relu)
```



```

In [4]: 1 np.random.seed(900)
        2 # read in dog image as np array
        3 img = np.array(imread("dog.jpg"))
        4
        5 fig, ax = plt.subplots(1, 4, figsize=(25, 5))
        6
        7 # generate 4 different random kernels to see different affects
        8 for i in range(4):
        9     # first layer
        10     # randomly generate a 3x3 kernel
        11     rand_kernel = np.random.randn(3,3)
        12
        13     # normalize the values to fall within [-0.5, 0.5]
        14     img_norm = normalize(img, -0.5, 0.5)
        15
        16     # apply 2D convolution
        17     img_conv = signal.convolve2d(img_norm, rand_kernel, mode='same')
        18
        19     # apply relu
        20     img_relu = np.clip(img_conv, 0, None)
        21
        22     # second layer
        23     # randomly generate a second 3x3 kernel
        24     rand_kernel2 = np.random.randn(3,3)
        25
        26     # normalize the values to fall within [-0.5, 0.5]
        27     img_norm2 = normalize(img_relu, -0.5, 0.5)
        28
        29     # apply second 2D convolution
        30     img_conv2 = signal.convolve2d(img_norm2, rand_kernel2, mode='same')
        31
        32     # apply relu
        33     img_relu2 = np.clip(img_conv2, 0, None)
        34     ax[i].imshow(img_relu2)

```



Answers

The convolved images accentuated the eyes, the nose and the edges of its ears of the dog while de-emphasized the overall figure. The effect of the second-layer convolution operation on the re-normalized filtered images detected the overall figure. Therefore we think the first few layers of a CNN detect some local features while the additional deeper layers detect more overall/abstract features.

Problem 2: Running a Convolutional Neural Network (20 pts)

Part A: Building the Model

In this first part, you will create a convolutional neural network using Keras to predict the type of object in an image. Load the [CIFAR-10 \(https://keras.io/datasets/#cifar10-small-image-classification\)](https://keras.io/datasets/#cifar10-small-image-classification) dataset, which contains 50,000 32x32 training images and 10,000 test images of the same size, with a total of 10 sizes.

Use a combination of the [following layers \(https://keras.io/layers/convolutional/\)](https://keras.io/layers/convolutional/): Conv2D, MaxPooling2D, Dense, Dropout and Flatten Layers (not necessarily in this order). You may use an existing architecture like AlexNet or VGG16, or create one of your own design. However, you should write your own layers and not use a pre-written implementation.

Convolutional neural networks are very computationally intensive. We highly recommend that you train your model on JupyterHub using GPUs. On CPUs, this training can take up to several hours. On GPUs, it can be done within minutes.

1. Report the total number of parameters.
2. How does the number of total parameters change (linearly, exponentially) as the number of filters per layer increases?
3. Generate a plot showing this relationship and explain why.

For instance, start by assigning 32 filters to each Conv2D layer, then 64, 128, etc. and recording the total number of parameters for each model.

```

In [5]: 1 def cnn_model(optimizer, n_filters, x_train):
2         # model
3         model = Sequential()
4         model.add(Conv2D(n_filters, (3, 3), padding='same',
5                           input_shape=x_train.shape[1:]))
6         model.add(Activation('relu'))
7         model.add(Conv2D(n_filters, (3, 3), padding='same',
8                           input_shape=x_train.shape[1:]))
9         model.add(Activation('relu'))
10        model.add(Conv2D(n_filters, (3, 3), padding='same',
11                          input_shape=x_train.shape[1:]))
12        model.add(Activation('relu'))
13        model.add(Conv2D(n_filters, (3, 3)))
14        model.add(Activation('relu'))
15        model.add(MaxPooling2D(pool_size=(2, 2)))
16        model.add(Dropout(0.25))
17
18        model.add(Conv2D(n_filters*2, (3, 3), padding='same'))
19        model.add(Activation('relu'))
20        model.add(Conv2D(n_filters*2, (3, 3), padding='same'))
21        model.add(Activation('relu'))
22        model.add(Conv2D(n_filters*2, (3, 3), padding='same'))
23        model.add(Activation('relu'))
24        model.add(Conv2D(n_filters*2, (3, 3)))
25        model.add(Activation('relu'))
26        model.add(MaxPooling2D(pool_size=(2, 2)))
27        model.add(Dropout(0.25))
28
29        model.add(Conv2D(n_filters*4, (3, 3), padding='same'))
30        model.add(Activation('relu'))
31        model.add(Conv2D(n_filters*4, (3, 3), padding='same'))
32        model.add(Activation('relu'))
33        model.add(Conv2D(n_filters*4, (3, 3), padding='same'))
34        model.add(Activation('relu'))
35        model.add(Conv2D(n_filters*4, (3, 3)))
36        model.add(Activation('relu'))
37
38        model.add(MaxPooling2D(pool_size=(2, 2)))
39        model.add(Dropout(0.25))
40
41        model.add(Flatten())
42        model.add(Dense(512))
43        model.add(Activation('relu'))
44        model.add(Dropout(0.5))
45        model.add(Dense(num_classes))
46        model.add(Activation('softmax'))
47
48        # compile the model using optimizer
49        model.compile(loss='categorical_crossentropy',
50                      optimizer=optimizer,
51                      metrics=['accuracy'])
52        return model

```

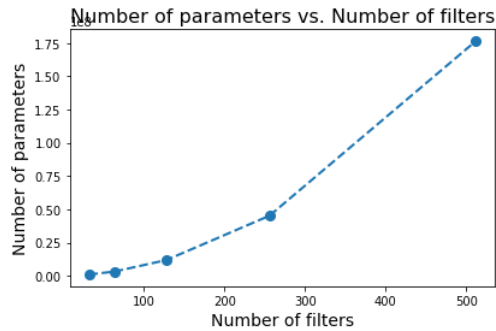
```
In [6]: 1 # read in data
2 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
3
4 # specify some model parameters
5 num_classes = len(np.unique(y_train))
6 batch_size = 32
7 epochs = 25
8 n_filters = 32
9 optimizer = keras.optimizers.rmsprop(lr=0.0001, decay=1e-6)
10 # sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)
11
12 # one hot encode the response y
13 y_train_onehot = keras.utils.to_categorical(y_train)
14 y_test_onehot = keras.utils.to_categorical(y_test)
15
16 model = cnn_model(optimizer, n_filters, x_train)
17 model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
activation_1 (Activation)	(None, 32, 32, 32)	0
conv2d_2 (Conv2D)	(None, 32, 32, 32)	9248
activation_2 (Activation)	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
activation_3 (Activation)	(None, 32, 32, 32)	0
conv2d_4 (Conv2D)	(None, 30, 30, 32)	9248
activation_4 (Activation)	(None, 30, 30, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
conv2d_5 (Conv2D)	(None, 15, 15, 64)	18496
activation_5 (Activation)	(None, 15, 15, 64)	0
conv2d_6 (Conv2D)	(None, 15, 15, 64)	36928
activation_6 (Activation)	(None, 15, 15, 64)	0
conv2d_7 (Conv2D)	(None, 15, 15, 64)	36928
activation_7 (Activation)	(None, 15, 15, 64)	0
conv2d_8 (Conv2D)	(None, 13, 13, 64)	36928
activation_8 (Activation)	(None, 13, 13, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_2 (Dropout)	(None, 6, 6, 64)	0
conv2d_9 (Conv2D)	(None, 6, 6, 128)	73856
activation_9 (Activation)	(None, 6, 6, 128)	0
conv2d_10 (Conv2D)	(None, 6, 6, 128)	147584
activation_10 (Activation)	(None, 6, 6, 128)	0
conv2d_11 (Conv2D)	(None, 6, 6, 128)	147584
activation_11 (Activation)	(None, 6, 6, 128)	0
conv2d_12 (Conv2D)	(None, 4, 4, 128)	147584
activation_12 (Activation)	(None, 4, 4, 128)	0
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 128)	0
dropout_3 (Dropout)	(None, 2, 2, 128)	0
flatten_1 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
activation_13 (Activation)	(None, 512)	0
dropout_4 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130
activation_14 (Activation)	(None, 10)	0
Total params: 942,314		
Trainable params: 942,314		

Non-trainable params: 0

```
In [7]: 1 # model = cnn_model(optimizer, n_filters=512)
        2 # model.summary()
```

```
In [8]: 1 n_params = [942314, 3224522, 11825546, 45174026, 176456714]
        2 n_filters = [32, 64, 128, 256, 512]
        3
        4 plt.plot(n_filters, n_params, 'o--', linewidth=2, markersize=8)
        5 plt.xlabel("Number of filters", fontsize=14)
        6 plt.ylabel("Number of parameters", fontsize=14)
        7 plt.title("Number of parameters vs. Number of filters", fontsize=16)
        8 plt.tight_layout()
        9 plt.show()
```



Answers

1. Report the total number of parameters.

As summarized above, our model has a total of 942,314 parameters.

2. How does the number of total parameters change (linearly, exponentially) as the number of filters per layer increases?

The total number of parameters within the entire network increases approximately linearly (with slightly increasing slope) as the number of filters grows.

3. Generate a plot showing this relationship and explain why.

This relationship is plotted above. This is because Convolutional Neural Networks share the same parameters across all kernels in the same layer. Suppose the biggest kernel size is $O(m^2)$, L layers in total, then total # parameters in our CNN =

$$\#Conv = n_{filters} (m^2 + 1) + 2^1 n_{filters} (m^2 + 1) + \dots + 2^{L-1} n_{filters} (m^2 + 1)$$

$$\#Dense = 2^{L-1} n_{filters} (m^2 + 1)(512 + 1) + (512)(10)$$

$$\#Total = \#Conv + \#Dense = O(m^2 2^L (n_{filters}))$$

Part B: Training and Evaluating the Model

Now train your model. You can choose to train your model for as long as you'd like, but you should aim for at least 10 epochs. Your validation accuracy should exceed 70%. Training for 10 epochs on a CPU should take about 30-60 minutes.

```
In [9]: 1 ## fit model
        2 # model.fit(x_train, y_train_onehot,
        3 #           batch_size=batch_size,
        4 #           epochs=epochs,
        5 #           validation_data=(x_test, y_test_onehot),
        6 #           shuffle=True)
        7
        8 ## Save model and weights
        9 # save_dir = os.path.join(os.getcwd(), 'saved_models')
        10 # model_name = 'keras_cifar10_trained_model3.h5'
        11
        12 # if not os.path.isdir(save_dir):
        13 #     os.makedirs(save_dir)
        14 # model_path = os.path.join(save_dir, model_name)
        15 # model.save(model_path)
        16 # print('Saved trained model at %s ' % model_path)
```

```
In [10]: 1 model = load_model('saved_models/keras_cifar10_trained_model3.h5')
```

```
In [11]: 1 scores = model.evaluate(x_test, y_test_onehot, verbose=1)
2         print('Test loss:', scores[0])
3         print('Test accuracy:', scores[1])
```

```
10000/10000 [=====] - 4s 357us/step
Test loss: 0.8020128004074096
Test accuracy: 0.758
```

Part C: Visualizing the Feature Maps

We would also like to examine the feature maps that are produced by the intermediate layers of the network.

Using your model, extract 9 feature maps from an intermediate convolutional layer of your choice and plot the images in a 3x3 grid. Also plot your original input image (choose an image of your choice).


You may use the helper function `get_feature_maps()` to extract weights.

```
In [12]: 1 def get_feature_maps(model, layer_id, input_image):
2         model_ = Model(inputs=[model.input], outputs=[model.layers[layer_id].output])
3         return model_.predict(np.expand_dims(input_image, axis=0))[0,:,:,:].transpose((2,0,1))
```

Original image


```
In [13]: 1 plt.imshow(x_train[4])
```

```
Out[13]: <matplotlib.image.AxesImage at 0x7fcea651ef0>
```



9 extracted feature maps

```
In [14]: 1 feature_maps = get_feature_maps(model, 3, x_train[4][:9])
2         fig, ax = plt.subplots(3, 3, figsize=(10, 10))
3         for i, feat_map in enumerate(feature_maps):
4             if i <= 2:
5                 ax[0][i].imshow(feat_map)
6             elif i <= 5:
7                 ax[1][i-3].imshow(feat_map)
8             else:
9                 ax[2][i-6].imshow(feat_map)
```



Problem 3: Recurrent Neural Networks (20 pts)

Learning to add numbers with a recurrent neural network

In this exercise, we will be using using recurrent neural network to add three digit numbers, encoded as character strings.

For example, given a string '223+12', we would like to return '235', without teaching the model explicit addition rules.

You are given the class **CharacterTable** to assist with encoding and decoding, which is initialized below:

```
In [15]: 1 from HW6_functions import *
2 chars = '0123456789+ '
3 ctable = CharacterTable(chars)
```

CharacterTable contains functions *encode* and *decode*.

encode takes in a string and the number of rows needed in the one hot encoding.

decode returns the string corresponding to the encoded one hot encoding.

An example of usage below:

```
In [16]: 1 encoded_123 = ctable.encode('123', 3)
2 print("Encoded Format: \n {}".format(encoded_123))
3 decoded_123 = ctable.decode(encoded_123)
4 print("Decoded Format: {}".format(decoded_123))
```

```
Encoded Format:
[[0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]]
Decoded Format: 123
```

Generating Training Data

Your first task is to create the data to train on. Luckily, we have virtually unlimited training data because addition is trivial for Python.

You will populate two arrays, *problems* and *answers*, which contain your predictors and target variables.

Examples from *problems*:

```
In [17]: 1 '    1+7 '
2
3 '   12+10 '
4
5 '520+880 '
```

```
Out[17]: '520+880 '
```

Examples from *answers*:

```
In [18]: 1 '8 '
2
3 '22 '
4
5 '1400 '
```

```
Out[18]: '1400 '
```

Notice that spaces are inserted to the left and right of strings within *problems* and *answers* to keep the dimensions of the input and output the same. When adding three digit numbers, the maximum possible length of a string in *problems* is 7, while the maximum possible length of a string in *answers* is 4.

```
In [19]: 1 # TRAINING_SIZE = 50000
2 # DIGITS = 3
3 # MAXLEN = DIGITS + 1 + DIGITS
```

```
In [20]: 1 # problems = []
2 # answers = []
```

1. Populate the two matrices *X* and *y*, which contain the encoded version of *problems* and *answers*. The *i* th row in both matrices should contain one encoded problem and answer, respectively.

2. Next, shuffle your data and split it into training and validation sets. These matrices should be named *x_train*, *y_train*, *x_val*, and *y_val*.

```
In [21]: 1 def generate_training_data(size, n_digits):
2         # max number is 10 to the power of n_digits
3         MAX_NUMBER = np.power(10, n_digits)
4
5         # set maximum length for problems and answers
6         MAXLEN_PROB = n_digits + 1 + n_digits
7         MAXLEN_ANS = n_digits + 1
8
9         # randomly generate 2 sets of numbers, each of size=size
10        first_number = np.random.randint(MAX_NUMBER, size=size)
11        second_number = np.random.randint(MAX_NUMBER, size=size)
12
13        # generate strings of problems
14        problems = [str(x) + '+' + str(y) for x, y in zip(first_number, second_number)]
15        answers = [str(x) for x in (first_number + second_number)]
16        # add spaces to in front of the problems and answers that are not at max lengths
17        for i, (prob, ans) in enumerate(zip(problems, answers)):
18            if len(prob) < MAXLEN_PROB:
19                problems[i] = ' ' * (MAXLEN_PROB-len(prob)) + prob
20            if len(ans) < MAXLEN_ANS:
21                answers[i] = ans + ' ' * (MAXLEN_ANS-len(ans))
22        return problems, answers
```

```
In [22]: 1 # generate problems and answers
2         TRAINING_SIZE = 50000
3         DIGITS = 3
4         problems, answers = generate_training_data(size=TRAINING_SIZE, n_digits=DIGITS)
```

```
In [23]: 1 # encode problems and answers
2         X = np.array([ctable.encode(x, 7) for x in problems])
3         y = np.array([ctable.encode(y, 4) for y in answers])
```

```
In [24]: 1 X.shape, y.shape
```

```
Out[24]: ((50000, 7, 12), (50000, 4, 12))
```

```
In [25]: 1 # split into training and validation sets
2         from sklearn.model_selection import train_test_split
3         X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
4         X_train.shape, X_val.shape, y_train.shape, y_val.shape
```

```
Out[25]: ((40000, 7, 12), (10000, 7, 12), (40000, 4, 12), (10000, 4, 12))
```

Building the Model

1. Using Keras, create a recurrent model that takes in X and returns y . You are free to choose LSTM, or just a vanilla RNN to implement your model. Your model should take in NUM_LAYERS as a parameter.

2. Create and train models with 1, 2, and 3 layers over 50 epochs. Plot test accuracy as a function of epoch for each model. Note: You do not have to print the progress bars for each model in your final report, you only have to include the accuracy plots.

3. Which model has the highest test accuracy? By looking at the accuracy over epochs, what can you say about how depth affects training and performance for recurrent models?

```
In [26]: 1 BATCH_SIZE = 128
2         LAYERS = [1, 2, 3]
3         EPOCHS = 50
4         NODES = 128
5         MAXLEN = DIGITS + 1 + DIGITS
6         CHAR_LEN = len(chars)
7         DIGITS = 3
```

```
In [27]: 1 def LSTM_model(n_nodes, n_layers, maxlen, chars_length, n_digits):
2         model = Sequential()
3         # first layer
4         model.add(LSTM(n_nodes, input_shape=(maxlen, chars_length)))
5         model.add(layers.RepeatVector(n_digits + 1))
6         # add layers based on LAYERS variable
7         for _ in range(n_layers):
8             model.add(LSTM(n_nodes, return_sequences=True))
9
10        # Apply a dense layer to the every temporal slice of an input. For each of step
11        # of the output sequence, decide which character should be chosen.
12        model.add(layers.TimeDistributed(layers.Dense(chars_length)))
13        model.add(layers.Activation('softmax'))
14        model.compile(loss='categorical_crossentropy',
15                      optimizer='adam',
16                      metrics=['accuracy'])
17        return model
```

```

In [28]: 1 val_loss_final = []
2 val_acc_final = []
3 val_loss_epoch = []
4 val_acc_epoch = []
5 for layer in LAYERS:
6     model = LSTM_model(NODES, layer, MAXLEN, CHAR_LEN, DIGITS)
7     print('Layer: ', layer)
8     print('Fitting model')
9     history = model.fit(X_train, y_train, epochs=EPOCHS, batch_size=BATCH_SIZE, validation_data=(X_val, y_val), verbose=0, shuffle=
10     val_loss_epoch.append(history.history['val_loss'])
11     val_acc_epoch.append(history.history['val_acc'])
12
13     # calculate final scores
14     print('Final Validation...')
15     scores = model.evaluate(X_val, y_val, verbose=1)
16     print('Validation loss:', scores[0])
17     print('Validation accuracy:', scores[1])
18     print('-----')
19     val_loss_final.append(scores[0])
20     val_acc_final.append(scores[1])
21
22     # Save model and weights
23     save_dir = os.path.join(os.getcwd(), 'saved_models')
24     model_name = 'LSTM_layer_' + str(layer) + '.h5'
25
26     if not os.path.isdir(save_dir):
27         os.makedirs(save_dir)
28     model_path = os.path.join(save_dir, model_name)
29     model.save(model_path)
30     print('Saved trained model at %s ' % model_path)
31

```

```

Layer: 1
Fitting model
Final Validation...
10000/10000 [=====] - 2s 172us/step
Validation loss: 0.034280417773127556
Validation accuracy: 0.990525
-----
Saved trained model at /home/ubuntu/AC209b/HW6/109/saved_models/LSTM_layer_1.h5
Layer: 2
Fitting model
Final Validation...
10000/10000 [=====] - 2s 219us/step
Validation loss: 0.01872033587358892
Validation accuracy: 0.99365
-----
Saved trained model at /home/ubuntu/AC209b/HW6/109/saved_models/LSTM_layer_2.h5
Layer: 3
Fitting model
Final Validation...
10000/10000 [=====] - 3s 270us/step
Validation loss: 0.09582543868124485
Validation accuracy: 0.9684
-----
Saved trained model at /home/ubuntu/AC209b/HW6/109/saved_models/LSTM_layer_3.h5


```

Plotting accuracies

```

In [29]: 1 fig, ax = plt.subplots(1, 3, figsize=(20, 5))
2 epoch = [j for j in range(1, EPOCHS+1)]
3 for i, acc in enumerate(val_acc_epoch):
4     ax[i].plot(epoch, acc, 'o--', linewidth=2, markersize=8)
5     ax[i].set_xlabel("Epoch Number", fontsize=14)
6     ax[i].set_ylabel("Validation Accuracy", fontsize=14)
7     ax[i].set_title("Validation Accuracy vs. Epoch \n # Layer = {}".format(i+1), fontsize=16)
8

```



```

In [ ]: 1

```

